

Ryan Timbrook

Data Science 450, Spring 2017

Date: 04/30/2017

Assignment 2

Description: Decisions Trees and Classification Targeted Marketing Campaign

In this problem we will use historical data from past customer responses to build a classification model. We will apply the trained model to a new set of prospects to whom we may want extend an offer for a PEP. Rather than doing a mass marketing campaign to all new prospects, we would like to target those that are likely to respond positively to our offer (according to our classification model). Use the dataset provided (comma delimited, and the first row contains the field names) - bank-data.csv

A. Open the .CSV file and inspect the attributes, is there one or more attributes that should be removed? If so, do it now.

B. Optional - You can choose to use the ARFF format or use CSV, - Save the file in .ARFF format
- Load the .ARFF formatted data file

C. Use Azure Machine Learning or a machine learning tool of your choice to build a classification model. You can choose to use any of the classification algorithms provided.

1. Evaluate your model accuracy

Spend time experimenting with different setting for the treebased algorithms - such as pruning, binary branching, etc. to see if you can improve the performance of your model.

Consider performing parameter sweeping to optimize for specific evaluation metrics.

3. Review the structure of the tree

See Table 1 below for tree representation of the models base structure using all parameters as a formula.

This base models accuracy score is 84% with an AUC of 91%

4. Identify the more important feature information.

See Table 3.1 below which represents the important feature information associated with its weighting

5. Generate the ROC curve for your final model.

See table 3.3 below for the final models AUC ROC curve

The final models accuracy score is 90.5% with an AUC of 91%

6. Show a screenshot of the final decision tree and model accuracy statistics obtained from your model.

See tables 3 below for final models accuracy statistics and AUC scores.

Accuracy: 90.5%

AUC: 91%

Table: 1.1 Single Tree plot – All base attributes




Table: 1.2

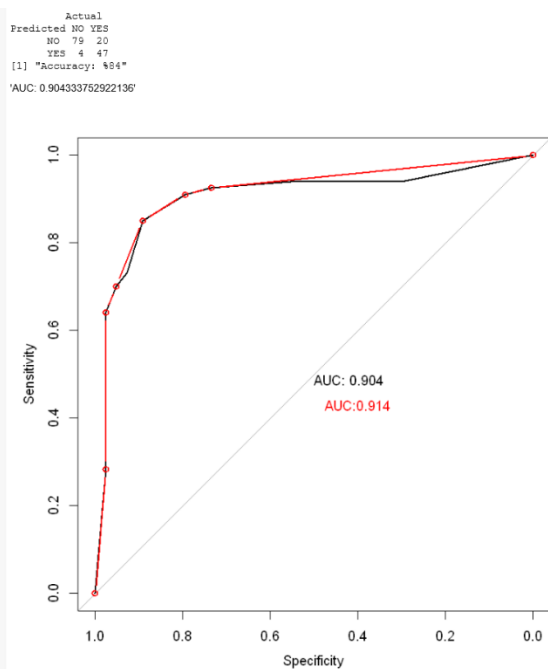


Table: 1.3




Table: 2.1 Experiment with Pruning

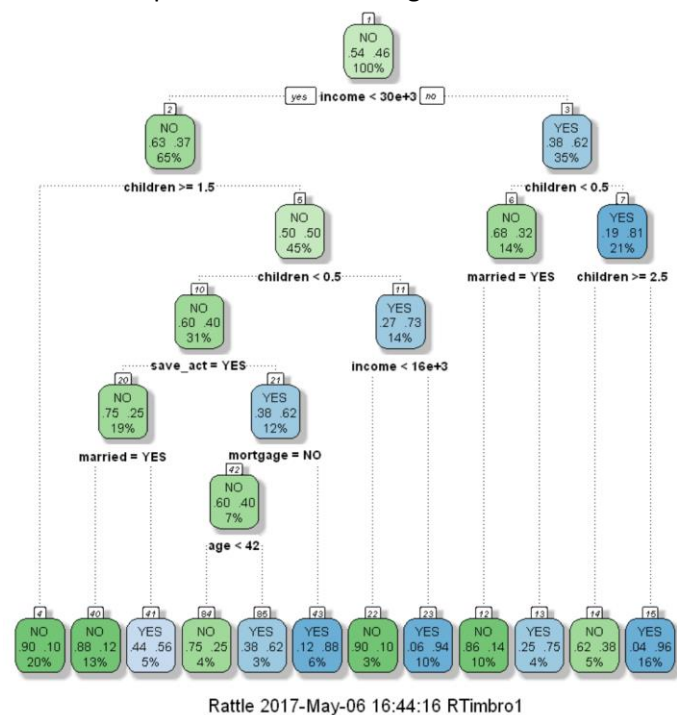


Table: 2.2

Actual
Predicted NO YES
NO 75 20
YES 6 47
(1) "Accuracy: 904"
"AUC: 0.90433762922136"




Table 2.3:

Root node error: 138/300 = 0.46

n= 300

	CP	nsplit	rel error	xerror	xstd
1	0.181159	0	1.00000	1.00000	0.062554
2	0.108696	1	0.81884	1.00725	0.062586
3	0.068841	2	0.71014	0.82609	0.060921
4	0.065217	4	0.57246	0.64493	0.057332
5	0.057971	5	0.50725	0.61594	0.056557
6	0.043478	6	0.44928	0.57246	0.055280
7	0.028986	7	0.40580	0.50000	0.052819
8	0.014493	9	0.34783	0.47101	0.051707
9	0.010000	11	0.31884	0.43478	0.050204




Table 3.1: Parameter Sweep model cp selection

CART

600 samples
10 predictor
2 classes: 'NO', 'YES'

No pre-processing
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 539, 541, 541, 541, 540, 540, ...
Resampling results across tuning parameters:

cp	Accuracy	Kappa
0.000000000	0.8801681	0.7565883
0.001824818	0.8868922	0.7702459
0.018248175	0.8836135	0.7636111
0.032846715	0.8266287	0.6483315
0.034671533	0.8199602	0.6344083
0.040145985	0.8019001	0.5972433
0.058394161	0.7717542	0.5367789
0.062043796	0.7667250	0.5269858
0.098540146	0.6963740	0.3666974
0.124087591	0.6271390	0.2036250

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was cp = 0.001824818.

income	70.4343294945446
children	54.178777167254
mortgageYES	28.8626761973173
save_actYES	24.6513080259642
marriedYES	22.0554065499718
age	21.4636585991491
regionRURAL	4.36771415189027
current_actYES	1.3824058287796
regionSUBURBAN	1.29657996357013
sexMALE	1.29657996357013
carYES	1.02050841274655

Final Model feature importance list

Table 3.2: Parameter Sweep, Final Model

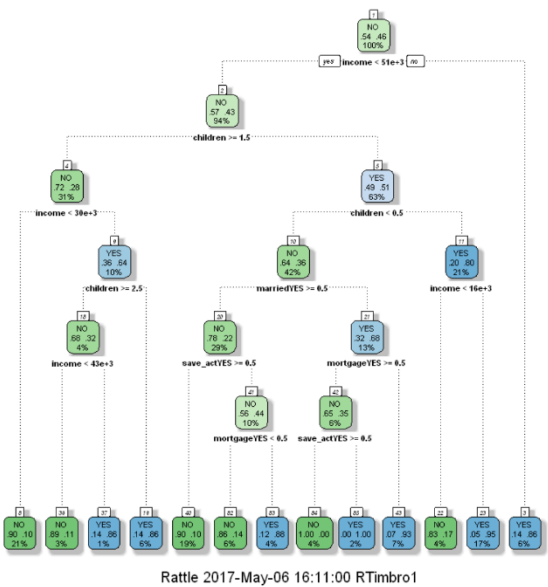


Table 3.3: Parameter Sweep – Final Model AUC

Actual
Predicted NO YES
NO 304 35
YES 22 239
[1] "Accuracy: 90.58"
'AUC: 0.920704401952443'

